

The Utility of NLPPhrases Using NFL Scouting Reports to Predict NFL Rookie Year Success

Sammy Muench Gabe Schwartz Jacqueline Kloner

January 19 2026

abstract

NFL teams invest heavily in draft evaluation, yet finding undervalued players beyond the early rounds remains difficult. Public scouting reports on NFL.com summarize film-based assessments in free text, but this information is rarely turned into quantitative features for prediction. We extract structured, interpretable trait scores from NFL.com scouting narratives (2014–2024) by prompting Google Gemini 1.5 Flash to map phrases from each report to a fixed set of football-relevant categories with integer sentiment scores. We then test whether these text-derived traits alone can predict rookie-year success, defined as an above-average Pro Football Focus (PFF) grade (≥ 60) with a minimum snap threshold. Using year-based splits (train: 2014–2020, validate: 2021, test: 2022–2024), we evaluate models by how well they rank players into a draft board, measuring the number of true successes among the top- K recommended prospects. While the model that only uses scouting reports is less predictive than the true draft, combining it with other signals preserves top- K hit rates while shifting hits later in the draft, suggesting that scouting-report text can help surface hidden gems. Overall, our results show that unstructured scouting narratives contain measurable signal about NFL readiness and can serve as a practical complement to existing draft evaluation tools.

1 Introduction

The NFL Draft is an annual event where professional American football teams pick players from college in inverse order of strength; that is, the weakest team chooses first from the pool of eligible college players. Despite months of preparation leading up to the draft, predicting long-term success translation from college to the NFL remains a major challenge. For example, only 27% of first-round wide receivers drafted between 2000 and 2019 signed a second contract with the team that drafted them. [9]

While teams rely on a combination of film study, NFL combine performance, and in-person interviews, these methods remain largely qualitative and are often subjective. Analysts and laypeople take to social media to divulge their opinions on college players; professionals create written scouting reports; in-house scouts offer counsel to their own teams [6]. Combining this rich pre-draft evaluation process with the NFL combine and pro days leads teams, analysts, and casual followers to some consensus on where players should be drafted, with some individual analysts achieving a 0.7 correlation between their predicted draft and the true draft [4]. The NFL directly contributes to public consensus on players by posting scouting reports on most players declared for the draft. Their consistent posting of these evaluations over the last ten years has resulted in a large set of unstructured text-based reports for every player. This trove of knowledge has historically gone unanalyzed pre-draft. If we can extract meaningful information from this text and link it to the draft, we may be able to improve many teams’ draft day decisions.

1.1 Previous Work: Predicting Rookie Performance

Efforts to predict rookie performance across various professional sports have produced mixed results, underscoring the inherent unpredictability of professional sports. Dedrick’s Predicting the Success of NFL Rookies: Production vs. Efficiency found that factors like age, speed, and college production

predicted outcomes for some positions (WRs, TEs) but not others (RBs). More importantly, the study found no significant correlation between college performance and NFL production, reinforcing the difficulty of translating amateur statistics to the pro level [2]. Similarly, Czaronis et al.’s How to Predict the Performance of NBA Draft Prospects introduced a relevance-based model that improved predictions but still leaned heavily on flawed college-team metrics [1]. These works expose the limitations of traditional quantitative inputs.

This gap has sparked growing interest in AI and NLP, which allow researchers to analyze unstructured text alongside numerical data. Kang’s Word Analysis Approach to NFL Draft Projection showed that linguistic traits in scouting reports may predict QB performance, identifying quick release and defensive reads as high-signal traits. This study’s goal aligns closely with ours, but Kang’s work has several important limitations that we address. Kang uses TF-IDF bigrams to categorize words into traits, which may leave out key context; for example, a report could say “solid pass blocking” or “elite pass blocking,” but bigrams would leave out the modified “solid” vs “elite.” Additionally, TF-IDF does not capture football knowledge, while our study uses Google Gemini Flash 1.5, a GPT trained on the broader web. This model has a built-in understanding of football and can represent scouting text with that knowledge in mind [5].

Kang noted that better players may have longer scouting reports because they are more likely to drive traffic to the NFL’s website for fans to read. This reality exists regardless of scouting report medium; better college players have more information written about them throughout the internet, so we move forward with knowledge of this potential bias that better players may have longer corpora. Using player traits is a key aspect of our study, one that has been explored before. Berri and Simmons (2011) evaluated the relationship between traits like height, BMI, and intelligence tests and quarterback success. Though these metrics showed some relationship to draft position, they lacked predictive signal relating to NFL quarterback performance [10].

Beyond performance prediction, NLP has revealed bias and systemic trends. Yellin’s Sidelined uncovered gender bias in NCAA basketball coverage, showing how language conveys value independent of actual performance [13]. Similarly, McManus’s Artificial Intelligence for Better In-Game NFL Performance highlighted AI’s role in decision-making and player safety. Studies like Juuri’s NFL game prediction work and Beal et al.’s soccer forecasting further confirm that combining structured and unstructured data yields better outcomes, though low-scoring variance and context-specific data remain challenges [7].

Across this literature, three themes emerge: (1) metrics like draft position or combine data explain part of the story but fall short of fully predicting player value; (2) unstructured text offers signal but is not used as efficiently as it could be; and (3) scouting language contains football-related terminology that should be more deeply accounted for. Our work directly builds on these insights, especially Kang’s study. Unlike prior studies constrained to narrow positions, limited sources, or short time frames, we analyze a decade of NFL.com scouting reports across multiple positions. We employ a large language model to translate freeform “Strengths” and “Weaknesses” into structured, explainable features. Most importantly, while past research shows that NLP can classify or augment structured stats, we demonstrate that it alone carries measurable predictive power, surfacing undervalued rookies in later rounds despite having no input from aggregate metrics. In doing so, we expand the scope of sports NLP from descriptive language analysis to actionable talent evaluation, laying groundwork for a richer literature on unstructured data in player forecasting.

1.2 Study Goals

Our research uses NFL scouting reports for three purposes: Player evaluation. Clear, succinct scouting metrics for each player may help NFL teams better understand player skill and fit. Using natural language processing, can we transform unstructured text-based scouting reports into meaningful quantitative features? Player prediction. If we are able to uncover meaningful qualitative features, do they contain any predictive signal on whether players will succeed in their rookie year? Model evaluation. When evaluating draft predictions, current literature tends to use traditional metrics like MAE or RMSE. [3] [8]. We argue that evaluating models with MAE and RMSE, metrics that treat each ob-

servation equally, or offer little incentive for teams to adopt machine learning models into their scouting departments. Teams spend millions of dollars to identify two primary sets of players: the best, and the undervalued. As such, models should be evaluated similarly. In keeping with that theme, we introduce a novel evaluation technique that involves ranking each model’s top choices and comparing those with true draft performance.

2 Methods

2.1 Data

We sourced predictors (X) from NFL.com and our response variable (Y) from Pro Football Focus (PFF). PFF’s proprietary grading system assigns each player, on every snap, a score from -2 to $+2$ in 0.5-point increments; these are aggregated into a 0–100 seasonal metric (with “average” around 60). A “snap” denotes each instance the center delivers the ball to the quarterback, so total snaps equate to the number of plays featuring the player.

Our dataset consisted of three blocks of information: (1) basic biographical and draft details—player name; height, weight, and wingspan (converted to decimal values); prospect grade; and overall pick (computed as $(\text{round}-1) \cdot 32 + \text{pick}$); (2) Next Gen Stats metrics—sectional scores and ranks cast to integers; and (3) pre-draft narrative text from the “Overview,” “Strengths,” and “Weaknesses” paragraphs. All features were consolidated into position-specific CSV files, with one row per player and one column per attribute. Not all positions were obtained due to data bottlenecks: QB, WR, and TE were not present in our final dataset, but all other football positions are present.

To convert dense scouting narratives into quantitative predictors, we addressed three primary challenges: scalability, domain-specific relevance, and output reliability. Manually converting tens of thousands of phrases across thousands of players to quantitative measures is infeasible. Bag-of-Words models, n-gram models, and SpaCy generated high-dimensional, context-insensitive feature vectors that were unhelpful for our task. Off-the-shelf sentiment tools misinterpret domain terminology, labeling “monstrous wide receiver” negative despite its positive connotation in football. To capture nuance, we used Google Gemini 1.5 Flash to classify each phrase into predefined performance categories (e.g. “Play Diagnosis and Instincts,” “Pursuit Angle”) with sentiment attribution. Google Gemini 1.5 Flash ingested each player’s scouting report with specific phrases. Our prompt included pre-determined categories for the LLM and methods for assigning point values for each metric. The LLM was given discretion to assign point magnitudes as it wished; most phrases had scores of -1 or 1 , but some fell as low as -5 or as high as 5 for especially strong language. Figure 1 contains the full data generating process, and Appendix 1 contains the full prompt given to Google Gemini 1.5 Flash and a sample output.

For each player, we sum the integer scores across all phrases in their report to obtain a final score per-category. Results were not deterministic, so to ensure that results made sense, we required Gemini to explain every decision made. Each feature is paired with an explanation of why Gemini scored the prompt the way it did, providing transparent, interpretable outputs. Appendix 2 highlights this.

We use Pro Football Focus (PFF) grades as our primary outcome variable because they represent the industry standard for player evaluation. All 32 NFL teams subscribe to PFF data, and the league itself regularly incorporates PFF grades into official player descriptions during broadcasts as a trusted measure of skill. Unlike coarse metrics such as Approximate Value or games played, which fail to capture the nuance of individual contributions, PFF grades are built from play-by-play aggregation and provide a uniquely granular view of performance.

The resulting pipeline generates position-specific feature sets with embedded rationales, ensuring that every score is traceable to explicit, LLM-provided reasoning. All ranking and top-k analyses in Section 5.2 are computed on the 2022–2024 test set only. Due to licensing restrictions, we cannot release raw PFF grades or the original scouting report text. We release derived trait scores for each anonymized player ID, which are sufficient to reproduce all analyses.

```

[list of all positions]
Loop 1: Loop through positions in list:
    [empty list of URLs]
    Loop 2: Loop through every year from 2014 to now:
        Go to URL, find number of pages for year/position combo
        Loop 3: Loop through each page of players:
            Add each player's URL to URL list
            Conclude Loop 3
        Conclude Loop 2
    Loop 4: Loop through all URLs in list:
        Go to URL
        Scrape necessary data
        Store in CSV labeled for current position
        Conclude Loop 4
    Conclude Loop 1

```

Figure 1: AI Algorithm for Trait Extraction

2.2 Evaluation

To transform a PFF grade into a binary classification variable, a threshold for a good grade must be chosen. Choosing this threshold leaves players with either a good grade or bad grade. Thus, all players with the same binary classification are treated equally, leading to elite players and slightly-above-mediocre players equally scored within the model. For this reason, we argue that simple binary classification metrics are not reasonable for this problem. Instead, for each model, we evaluate the probability of a player receiving a good grade given our model and data. We then “rank” players according to their predicted probability of success, creating a quasi-draft board. Finally, we calculate a cumulative sum of successful players as we move down each draft board. This simulation mirrors how NFL teams draft: given a master ranking of players, pick players according to the draft board.

Ranking Procedure. Let $\mathcal{D}_{\text{test}} = \{x_i\}_{i=1}^n$ denote the test set, where $x_i \in R^p$ is the feature vector for player i . Let $\hat{f}(\cdot)$ denote a trained logistic regression model.

For each test-set player i , we compute the predicted probability of success:

$$\hat{p}_i = \hat{f}(x_i) = \Pr(Y = 1 \mid X = x_i). \quad (1)$$

We define a ranking as a permutation of $\{1, \dots, n\}$ such that:

$$\hat{p}_{(1)} \geq \hat{p}_{(2)} \geq \dots \geq \hat{p}_{(n)}. \quad (2)$$

The ranked test-set list (i.e., the model’s draft board) is then given by:

$$\text{rank}(\mathcal{D}_{\text{test}}) = [(1), (2), \dots, (n)]. \quad (3)$$

3 Experiment

3.1 Exploratory Data Analysis

Position-specific and position-agnostic experiments were conducted. The below exploratory data analysis includes graphs from our position agnostic experiment. Although filtering data by position should isolate attributes that some positions may value more than others, doing so caused issues in dataset size. We offer remedies for this issue in Conclusions.

We observe correlations between most of our features that fall in line with NFL logic. Features with negative correlations are disjoint; for example, technical skills and size/body composition as




Figure 2: Relationship between speed and linear quickness.

skills usually must be improved separately, in the form of film study versus weight room training. Features with a positive correlation can be improved together; for example, technical fundamentals and instincts/play recognition are both improved with years of experience. Straight line speed and change of direction are often associated with lower-half athleticism and go hand-in-hand. Our features demonstrate this signal across most pairwise combinations, as shown in Table 1.

Table 1: Correlation matrix of attributes.

	Size & Body	Explosive Power	Speed	Technical	Instincts	E ort	Versatility	Durability	Agility
Size & Body Composition	1.000	-0.033	-0.212	-0.266	-0.256	-0.044	-0.009	-0.055	-0.287
Explosive Power	-0.033	1.000	0.472	-0.110	-0.168	0.137	0.081	-0.094	0.253
Speed & Linear Quickness	-0.212	0.472	1.000	-0.126	-0.150	-0.080	0.194	-0.080	0.611
Technical Fundamentals	-0.266	-0.110	-0.126	1.000	0.374	0.256	-0.046	0.058	0.138
Instincts & Play Recognition	-0.256	-0.168	-0.150	0.374	1.000	0.106	-0.037	0.046	0.049
E ort & Finishing Ability	-0.044	0.137	-0.080	0.256	0.106	1.000	-0.007	0.051	-0.038
Positional Versatility	-0.009	0.081	0.194	-0.046	-0.037	-0.007	1.000	0.010	0.152
Durability & Availability	-0.055	-0.094	-0.080	0.058	0.046	0.051	0.010	1.000	-0.060
Agility & Change of Direction	-0.287	0.253	0.611	0.138	0.049	-0.038	0.152	-0.060	1.000

3 2 Candidate Models

We are also interested in using these meaningful features to predict performance. We limited our response to rookie year performance. This decision comes with its own biases, namely that players do not always peak in their rookie year. We made this choice for simplicity. While it would be valuable to examine whether rookie-season performance is indicative of long-term outcomes, our study does not evaluate this relationship.

We reframed the problem as binary classification to distinguish likely successes from non-successes. We prioritized precision over recall, since falsely recommending a poor player imposes greater cost on




Figure 3: Play Recognition vs Technical Skills for all players. We observe a moderate positive correlation. We hypothesize that NFL players who have strong play recognition also have strong technical skills because both attributes relate to playing experience.

a team than overlooking an eventual high-performer. “Success” was defined as a rookie PFF grade ≥ 60 (above the median). To remove noise from limited playing time, any player with ≤ 100 snaps was labeled negative.

After excluding records lacking pre-draft data, position-specific datasets contained 150–400 observations and an all-position set contained 2,300 entries. We experimented with (1) random splits (fixed seed) and (2) temporal splits—training on 2014–2020, validating on 2021, and testing on 2022–2024. Model selection employed 5-fold cross-validation. Where relevant, we used GroupKFold by draft year to enforce temporal separation and account for evolving positional skill requirements. We compared Logistic Regression, tuning the regularization parameter C , and Random Forest models, tuning tree count and maximum depth.

Traditional error metrics assume that all test data points should be treated equally during evaluation. In professional sports’ drafting context, this is never true: a false positive, or a player predicted to succeed, has more ramifications than a false negative, where a good player is overlooked. In fact, sports media often measures drafting by “successes” and “busts,” which, in a modeling context, is the ratio of true positives to false positives, or precision. [11] [12] As such, when using draft prediction models to predict PFF grade or any continuous numeric grade, transforming the target variable to a binary “success” or “bust” reflects real-world incentives.

Final testing on heldout test sets relied on our novel evaluation method, described in section 5.

4 Results

4.1 Raw Results

Figure 4 shows all model results. The red curve is the baseline—as if we randomly selected each player. This graph works as a draft board, where we order each model’s player ranking by likelihood of success, then iteratively measure how often we are selecting good players as we move down the draft board. The blue line illustrates the “perfect draft” baseline of selecting all good players before the worst ones. The orange line contains the true drafts from 2022–2024: we order by draft position and iteratively take




Figure 4: Each curve shows the cumulative number of rookies with a good PFF grade (“successes”) as we move down the draft board, sorted by predicted probability from each model. A more concave curve indicates that the model is better at ranking true successes toward the top, just as an efficient draft board would prioritize the right players earlier. The red line is the random baseline, the blue line represents a hypothetical “perfect draft,” and the orange line reflects the actual NFL draft order. The light green curve represents a model trained on Next Gen Stats grades, and the light blue curve is our NLP model. Note that in the first 50 picks, the NLP model is comparable with NGS and the draft.

the cumulative sum of how many “good players” we have chosen at any given point. As such, a more concave curve means that our model is better at picking. This paper operates under the assumption that the true draft is a realistic proxy for the true landscape of how all 32 NFL teams believe players will perform.

The results show that the true NFL draft only slightly outperforms Next Gen Stats. This finding is consistent with the draft’s moderate correlation (0.574) with Next Gen Stats: players rated highly by NGS tend to be drafted earlier. However, the true draft still falls well short of the theoretical “perfect draft,” underscoring the difficulty teams face in identifying talent.

We observe that our model is worse than the others after about 50 players drafted. This result aligns with our expectations, as the only inputs to our model are three paragraphs about one scout’s view of a player. However, our model can lag behind the others in overall accuracy and still provide utility, especially if it identifies talent that may be overlooked by the others.

NFL teams use an ensemble of tools to build their draft boards: analytics teams, film, scouting teams, coach feedback, the NFL combine, and more. It is unreasonable to assume that a single statistical model can replace an NFL scouting department. In this vein, we evaluate our model by how it adds value to the current landscape of NFL draft scouting. We ensemble models together by taking the average of each model’s player rank. In Figure, we compare results from an ensemble of our baselines (NGS + draft) with the results of an ensemble of our baselines plus our model.

When we examined the intersection of all three predictive systems (the top-50 players for the true draft baseline, Next Gen Stats evaluations, and our model) we identified seven players whom all models classified as likely successes prior to the draft. Of these seven players, six produced successful rookie seasons according to our outcome metric. In other words, when a player was drafted highly, rated highly during the pre-draft evaluation process, and also received a strong projection from our model, this composite agreement was associated with a 6-out-of-7 rate of rookie-year success. For comparison, all first-round draft picks in our heldout set had a success rate of 23-out-of-39 (58%). Despite the low sample size, this result demonstrates that ensembling our NLP model together with

Table 2: Marginal gain of adding the NLP model to the NGS + NFL model. Similar to the Figure 3, this shows the cumulative # of good players selected as we move down our draft board. Though the two models have similar good player counts, we see a significant difference in the average draft position identified: the all-model ensemble identifies players with consistently later draft position with no loss in good player count. This result demonstrates that the NLP model can extract players for “cheaper,” or at later draft positions.

Top-	Hits ALL)	Hits NGS+Draft)	Avg Draft Pos ALL)	Avg Draft Pos NGS+Draft)	Hits	Draft Pos
10	8	8	20	7	0	13
20	16	15	29	15	1	14
50	30	35	55	30	-5	25
75	47	46	71	42	1	29
100	61	58	80	56	3	24
150	84	85	98	87	-1	11

other models may minimize risk in choosing the top players in the draft. Below is a case study of three players in our modeling framework, including Cooper DeJean, one player in the intersection of all models.

1. **Evan Neal:** Rated high by draft (5th overall in 3-year test set) but low by NGS (93rd) and NLP (193rd). Went on to bust.
2. **Kamren Kinchens:** hidden gem. Rated 17th by NLP model and 27th by NGS but was a late 3rd round pick. Now a sub-elite starter for the Rams (PFF).
3. **Cooper DeJean:** Ranked 10th by our model, 15th by NGS, 43rd in draft. Now an elite starter. These types of cases are our most successful - players who are moderately-to-highly ranked by all models tend to perform very well in the NFL.

5 Discussion

Overall, results reveal that the NFL scouting reports contain meaningful, predictive signal. Our model’s true positives largely consist of quality players who earned moderate playing time, averaging about half their teams’ snaps (438 total). These kinds of reliable role players are especially valuable in the later rounds of the draft (citation). Notably, the model also identified two undrafted players who passed our criteria for a good grade (assigned a draft slot of 300 for analysis). When NLP predictions are combined with existing information, such as draft position or Next Gen Stats, the overlap consistently produces higher-value picks. This reinforces the idea that aggregating models is the most effective approach, with NLP adding a distinct and complementary layer. Given that even the first round produced only a 58% hit rate in our heldout set, the ability to surface overlooked contributors highlights the practical value of this method.

While our NLP-derived metric does not outperform draft position or Next Gen Stats, both of which consolidate vast amounts of quantitative and qualitative scouting resources, it nevertheless demonstrates statistically meaningful signal given its simplicity. Among drafted players, the top 50 prospects ranked by our NLP model achieved a 62% success rate, compared to 84% for Next Gen Stats and 82% for draft position. Importantly, our approach relies solely on pre-draft scouting text, with no embedded information from aggregate performance metrics or draft capital. That it can achieve even modest predictive power relative to professional-grade metrics like Next Gen Stats underscores its value as a complementary lens. The contribution of our method lies not in replacing existing scouting systems but in showing that unstructured narrative reports contain hidden, quantifiable information about player potential that is not fully captured in traditional evaluation pipelines. We show that ensembling model selections together can minimize risk in drafting and uncover hidden gems later in the draft.

6 Limitations and Future Directions

Our position-specific models were hampered by data scarcity. After applying filters and cleaning, some heldout datasets contained as few as ten players. As a result, the models and data from this work was exclusively from the position-agnostic model which included all NFL players.

Although we initially anticipated that a deep neural network would outperform linear methods, it failed to do so. This issue likely stems from insufficient observations to support the MLP’s complexity. Future work should prioritize dataset expansion, either via additional seasons or alternative data sources, and further exploration of neural networks.

This problem could also be viewed as a 0-1 knapsack problem, where weights are a function of draft position and the values are the PFF grades. This type of study may more closely mirror the resource-allocation constraints of the NFL draft.

With a larger corpus, more positions to analyze, and more players per dataset, studies that build on these methods have the potential to unearth new scouting strategies that aggregate opinions across many scouts, professionals, and even laypeople to build more representative prediction models in the future. Our anonymized dataset and notebooks can be found on github.

Appendix

A Prompt

Below lists the prompt fed to Google Gemini:

Begin Prompt

You are given a blurb describing a prospective NFL player, of any position. Your task is to analyze the blurb in detail and categorize each meaningful phrase into the following pre-defined categories. After categorizing each phrase, provide a dictionary output EXACTLY in the form described below, with category names as keys and integer scores as values.

Rules for Scoring: The score for each category is calculated by counting each phrase that positively or negatively associates with the category. Phrases may be associated with more than one category if it covers more than one. Positive phrases count as +1 each. Negative phrases count as -1 each. The final score for each category is the sum of positive and negative phrase counts. If a phrase is especially glowing or poor, you may edit the score accordingly.

Size and Body Composition Height, weight, length and muscle mass relative to positional norms.

Explosive Power Ability to generate force in a single movement (first-step burst, punch, drive).

Speed and Linear Quickness Acceleration and top-end speed in straight-line sprints and pursuit.

Agility and Change-of-Direction Multi-plane footwork, hip fluidity and ability to transition sharply.

Technical Fundamentals Core position-specific skills (footwork, hand placement, throwing/tackling/route-running mechanics).

Instincts and Play Recognition Pre-snap alignment, post-snap keys, pattern recognition and anticipation. How rapidly a player reads the field, digests information, and makes the right call under game speed.

Effort and Finishing Ability Persistence through the whistle: sustaining blocks, driving through tackles, completing routes or throws.

Assignment Discipline and Reliability Consistency in sticking to scheme responsibilities, minimizing mental errors and blown assignments.

Positional Versatility Ability to line up and contribute at multiple spots or in varied roles/schemes.

Mental Resilience and Clutchness Capacity to stay calm under pressure, recover quickly from mistakes or adversity, and maintain focus.

Durability and Availability Injury history, recovery, endurance and consistent availability at full speed.

Leadership: Phrases specifically related to mental fortitude, loved by teammates, pregame speeches, leadership, trusted by teammates and coaches, dependable.

Instructions for output: Evaluate EVERY phrase in the blurb individually and categorize according to above definitions. Provide a SINGLE dictionary EXACTLY as formatted below, without any additional explanation or Python code around it. Use EXACTLY the category names listed above as keys. EVERY category must have a numeric score (0 if no phrases apply).

Example of desired sample dictionary output: {"Size and Body Composition": 3, "Explosive Power": -2, "Speed and Linear Quickness": 4, "Agility and Change-of-Direction": 1, "Technical Fundamen-

tals": -3, "Instincts and Play Recognition": 5, "Pad Level, Leverage and Balance": 2, "Effort and Finishing Ability": -4, "Assignment Discipline and Reliability": -4, "Positional Versatility": 0, "Mental Resilience and Clutchness": 9, "Durability and Availability": -1, "Leadership": 7}

After you write the dictionary EXACTLY in the format listed above, please list your justification after for each phrase. Your justification should be listed after the close bracket of the dictionary.

Here is the blurb: [Insert Player Description]

After receiving the prompt the LLM may respond in the following way:

{"Play Diagnosis and Instincts": -2,

...

"Competitive Motor and Effort": -1 }

****Justification:****

* **"Undersized outside linebacker lacking both starting experience and instincts for the position."**: This phrase negatively impacts "Play Diagnosis and Instincts". Lack of instincts is directly related to this category.

* **"He has a natural feel for navigating traffic, but that won't be very valuable to him until he improves his play recognition."**: This phrase is a positive attribute related to "Pursuit Speed and Angle Play" (navigating traffic implies good pursuit and angles), but the caveat about needing to improve play recognition negatively affects "Play Diagnosis and Instincts."

References

- [1] Megan Czasonis, Mark Kritzman, Cel Kulasekaran, and David Turkington. How to predict the performance of nba draft prospects. *SSRN Electronic Journal*, 01 2023.
- [2] Kylie Dedrick. Predicting the success of nfl rookies: Production vs. efficiency. *Renee Crown University Honors Thesis Projects, Syracuse University*, 2023.
- [3] Akshaj Enaganti and George Pappas. Optimizing nfl draft selections with machine learning classification. *AI*, 6(9):221, 2025.
- [4] Hogs Haven. Which analysts are best at predicting nfl draft outcomes? *Hogs Haven*, 2022.
- [5] David Kang. Word analysis approach to nfl draft projection and its impact on team performance. *Caroline Digital Repository*, 2017.
- [6] National Football League. College advisory committee. *NFL Football Operations*, 2024.
- [7] Christopher McManus. Artificial intelligence for better in-game nfl performance. *Honors Theses, Lee Honors College, Western Michigan University*, 2025.
- [8] Jason Mulholland and Shane T. Jensen. Predicting the draft and career success of tight ends in the national football league. *Journal of Quantitative Analysis in Sports*, 10(4):381–396, 2014.
- [9] Je Risdon. First-round draft hit rates: Tracking rates from the best (centers) to the worst (wrs). *Yahoo Sports*, 2024.
- [10] Rob Simmons and David Berri. Mixing the princes and the paupers: Pay and performance in the national basketball association. *Labour Economics*, 18(3):381–388, 2011.
- [11] Brent Sobleski. What are biggest hits and misses from 2025 nfl draft class so far? *Bleacher Report*, January 2025. Accessed: 2026-01-19.
- [12] Matt Verderame. The biggest nfl draft bust for every team. *Sports Illustrated*, 2023. Accessed: 2026-01-19.
- [13] Nathaniel Yellin. Sidelined.